

0. Install and Import Dependencies

```
!pip list

!pip install opencv-python matplotlib imageio gdown tensorflow

import os
import cv2
import tensorflow as tf
import numpy as np
from typing import List
from matplotlib import pyplot as plt
import imageio

tf.config.list_physical_devices('GPU')

physical_devices = tf.config.list_physical_devices('GPU')
try:
    tf.config.experimental.set_memory_growth(physical_devices[0],
True)
except:
    pass
```

1. Build Data Loading Functions

```
import gdown

url = 'https://drive.google.com/uc?id=1YlvpDLix3S-U8fd-gqRwPcWXAm8JwjL'
output = 'data.zip'
gdown.download(url, output, quiet=False)
gdown.extractall('data.zip')

def load_video(path:str) -> List[float]:
    cap = cv2.VideoCapture(path)
    frames = []
    for _ in range(int(cap.get(cv2.CAP_PROP_FRAME_COUNT))):
        ret, frame = cap.read()
        frame = tf.image.rgb_to_grayscale(frame)
        frames.append(frame[190:236, 80:220, :])
    cap.release()

    mean = tf.math.reduce_mean(frames)
    std = tf.math.reduce_std(tf.cast(frames, tf.float32))
    return tf.cast((frames - mean), tf.float32) / std
```

```

vocab = [x for x in "abcdefghijklmnopqrstuvwxyz'?!123456789 "]

char_to_num = tf.keras.layers.StringLookup(vocabulary=vocab,
oov_token="")
num_to_char = tf.keras.layers.StringLookup(
    vocabulary=char_to_num.get_vocabulary(), oov_token="", invert=True
)

print(
    f"The vocabulary is: {char_to_num.get_vocabulary()}"
    f"(size ={char_to_num.vocabulary_size()})"
)
char_to_num.get_vocabulary()
char_to_num(['n','i','c','k'])
num_to_char([14, 9, 3, 11])

def load_alignments(path:str) -> List[str]:
    with open(path, 'r') as f:
        lines = f.readlines()
    tokens = []
    for line in lines:
        line = line.split()
        if line[2] != 'sil':
            tokens = [*tokens, ' ',line[2]]
    return char_to_num(tf.reshape(tf.strings.unicode_split(tokens,
input_encoding='UTF-8'), (-1)))[1:]

def load_data(path: str):
    path = bytes.decode(path.numpy())
    #file_name = path.split('/')[-1].split('.')[0]
    # File name splitting for windows
    file_name = path.split('\\')[-1].split('.')[0]
    video_path = os.path.join('data','s1',f'{file_name}.mpg')
    alignment_path =
os.path.join('data','alignments','s1',f'{file_name}.align')
    frames = load_video(video_path)
    alignments = load_alignments(alignment_path)

    return frames, alignments

test_path = '.\\data\\s1\\bbal6n.mpg'

tf.convert_to_tensor(test_path).numpy().decode('utf-8').split('\\')[-1].split('.')[0]

frames, alignments = load_data(tf.convert_to_tensor(test_path))
plt.imshow(frames[40])

```

```

alignments

tf.strings.reduce_join([bytes.decode(x) for x in
num_to_char(alignments.numpy()).numpy()])

def mappable_function(path:str) ->List[str]:
    result = tf.py_function(load_data, [path], (tf.float32, tf.int64))
    return result

```

2. Create Data Pipeline

```

from matplotlib import pyplot as plt

data = tf.data.Dataset.list_files('./data/s1/*.mpg')
data = data.shuffle(500, reshuffle_each_iteration=False)
data = data.map(mappable_function)
data = data.padded_batch(2, padded_shapes=([75, None, None, None], [40]))
data = data.prefetch(tf.data.AUTOTUNE)
# Added for split
train = data.take(450)
test = data.skip(450)

len(test)

frames, alignments = data.as_numpy_iterator().next()

len(frames)

sample = data.as_numpy_iterator()

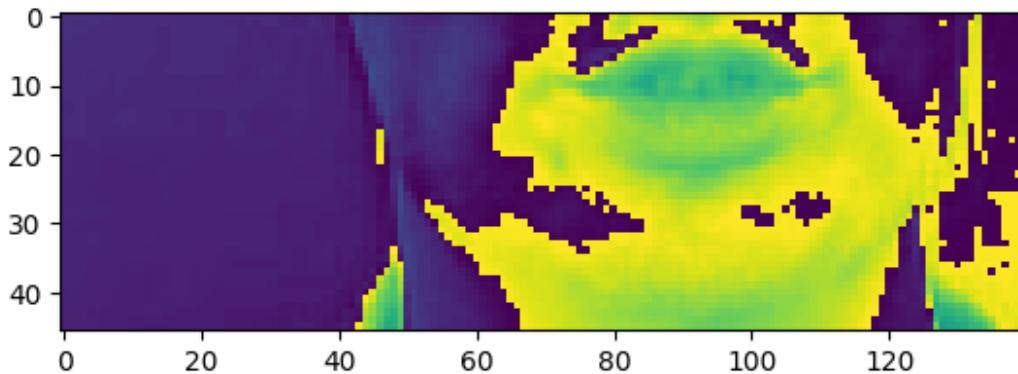
val = sample.next(); val[0]

imageio.mimsave('./animation.gif', val[0][0], fps=10)

# 0:videos, 0: 1st video out of the batch, 0: return the first frame
# in the video
plt.imshow(val[0][0][35])

<matplotlib.image.AxesImage at 0x10c2ead8d30>

```



```
tf.strings.reduce_join([num_to_char(word) for word in val[1][0]])
<tf.Tensor: shape=(), dtype=string, numpy=b'lay blue by e two please'>
```

3. Design the Deep Neural Network

```
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv3D, LSTM, Dense, Dropout,
Bidirectional, MaxPool3D, Activation, Reshape, SpatialDropout3D,
BatchNormalization, TimeDistributed, Flatten
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import ModelCheckpoint,
LearningRateScheduler

data.as_numpy_iterator().next()[0][0].shape
(75, 46, 140, 1)

model = Sequential()
model.add(Conv3D(128, 3, input_shape=(75,46,140,1), padding='same'))
model.add(Activation('relu'))
model.add(MaxPool3D((1,2,2)))

model.add(Conv3D(256, 3, padding='same'))
model.add(Activation('relu'))
model.add(MaxPool3D((1,2,2)))

model.add(Conv3D(75, 3, padding='same'))
model.add(Activation('relu'))
model.add(MaxPool3D((1,2,2)))

model.add(TimeDistributed(Flatten()))

model.add(Bidirectional(LSTM(128, kernel_initializer='Orthogonal',
return_sequences=True)))
model.add(Dropout(.5))
```

```

model.add(Bidirectional(LSTM(128, kernel_initializer='Orthogonal',
return_sequences=True)))
model.add(Dropout(.5))

model.add(Dense(char_to_num.vocabulary_size()+1,
kernel_initializer='he_normal', activation='softmax'))

model.summary()

```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv3d (Conv3D)	(None, 75, 46, 140, 128)	3584
activation (Activation)	(None, 75, 46, 140, 128)	0
max_pooling3d (MaxPooling3D)	(None, 75, 23, 70, 128)	0
conv3d_1 (Conv3D)	(None, 75, 23, 70, 256)	884992
activation_1 (Activation)	(None, 75, 23, 70, 256)	0
max_pooling3d_1 (MaxPooling 3D)	(None, 75, 11, 35, 256)	0
conv3d_2 (Conv3D)	(None, 75, 11, 35, 75)	518475
activation_2 (Activation)	(None, 75, 11, 35, 75)	0
max_pooling3d_2 (MaxPooling 3D)	(None, 75, 5, 17, 75)	0
time_distributed (TimeDistr ibuted)	(None, 75, 6375)	0
bidirectional (Bidirectiona l)	(None, 75, 256)	6660096
dropout (Dropout)	(None, 75, 256)	0
bidirectional_1 (Bidirectio nal)	(None, 75, 256)	394240
dropout_1 (Dropout)	(None, 75, 256)	0
dense (Dense)	(None, 75, 41)	10537

```
Total params: 8,471,924  
Trainable params: 8,471,924  
Non-trainable params: 0
```

$$5 * 17 * 75$$

```
yhat = model.predict(val[0])

1/1 [=====] - 3s 3s/step

tf.strings.reduce_join([num_to_char(x) for x in
tf.argmax(yhat[0],axis=1)])

<tf.Tensor: shape=(), dtype=string,
numpy=b'444iiiiiiiiiiimmmmmmmmmmmmmmmmmmmiiiiiiiiimmmmmmmmmiiiiiiiixxxxxx'>

tf.strings.reduce_join([num_to_char(tf.argmax(x)) for x in yhat[0]])

<tf.Tensor: shape=(), dtype=string,
numpy=b'444iiiiiiiiiiimmmmmmmmmmmmmmmmmmmiiiiiiiiimmmmmmmmmiiiiiiiixxxxxx'>

model.input_shape

(None, 75, 46, 140, 1)

model.output_shape

(None, 75, 41)
```

4. Setup Training Options and Train

```

        loss = tf.keras.backend.ctc_batch_cost(y_true, y_pred,
input_length, label_length)
    return loss

class ProduceExample(tf.keras.callbacks.Callback):
    def __init__(self, dataset) -> None:
        self.dataset = dataset.as_numpy_iterator()

    def on_epoch_end(self, epoch, logs=None) -> None:
        data = self.dataset.next()
        yhat = self.model.predict(data[0])
        decoded = tf.keras.backend.ctc_decode(yhat, [75,75],
greedy=False)[0][0].numpy()
        for x in range(len(yhat)):
            print('Original: ',
tf.strings.reduce_join(num_to_char(data[1][x])).numpy().decode('utf-
8'))
            print('Prediction: ',
tf.strings.reduce_join(num_to_char(decoded[x])).numpy().decode('utf-
8'))
            print('*'*100)

model.compile(optimizer=Adam(learning_rate=0.0001), loss=CTCLoss)

checkpoint_callback =
ModelCheckpoint(os.path.join('models','checkpoint'), monitor='loss',
save_weights_only=True)

schedule_callback = LearningRateScheduler(scheduler)

example_callback = ProduceExample(test)

model.fit(train, validation_data=test, epochs=100,
callbacks=[checkpoint_callback, schedule_callback, example_callback])

Epoch 1/100
2/450 [.....] - ETA: 3:03 - loss: 213.9969

```

KeyboardInterrupt

5. Make a Prediction

```

url = 'https://drive.google.com/uc?id=1vWscXs4Vt0a_1IH1-ct2TCgXAZT-
N3_Y'
output = 'checkpoints.zip'
gdown.download(url, output, quiet=False)
gdown.extractall('checkpoints.zip', 'models')

```

```

model.load_weights('models/checkpoint')

<tensorflow.python.checkpoint.checkpoint.CheckpointLoadStatus at
0x10cfb56c6a0>

test_data = test.as_numpy_iterator()

sample = test_data.next()

yhat = model.predict(sample[0])

1/1 [=====] - 1s 973ms/step

print('*'*100, 'REAL TEXT')
[tf.strings.reduce_join([num_to_char(word) for word in sentence]) for
sentence in sample[1]]

~~~~~ REAL TEXT

[<tf.Tensor: shape=(), dtype=string, numpy=b'place white at x six
please'>,
 <tf.Tensor: shape=(), dtype=string, numpy=b'lay blue in x four now'>]

decoded = tf.keras.backend.ctc_decode(yhat, input_length=[75,75],
greedy=True)[0][0].numpy()

print('*'*100, 'PREDICTIONS')
[tf.strings.reduce_join([num_to_char(word) for word in sentence]) for
sentence in decoded]

~~~~~ PREDICTIONS

[<tf.Tensor: shape=(), dtype=string, numpy=b'place white at x six
please'>,
 <tf.Tensor: shape=(), dtype=string, numpy=b'lay blue in x four now'>]

```

Test on a Video

```

sample = load_data(tf.convert_to_tensor('.\\data\\s1\\bras9a.mpg'))

print('*'*100, 'REAL TEXT')
[tf.strings.reduce_join([num_to_char(word) for word in sentence]) for
sentence in [sample[1]]]

~~~~~ REAL TEXT

[<tf.Tensor: shape=(), dtype=string, numpy=b'bin red at s nine
again'>]

```

```
yhat = model.predict(tf.expand_dims(sample[0], axis=0))

1/1 [=====] - 1s 720ms/step

decoded = tf.keras.backend.ctc_decode(yhat, input_length=[75],
greedy=True)[0][0].numpy()

print('~'*100, 'PREDICTIONS')
[tf.strings.reduce_join([num_to_char(word) for word in sentence]) for
sentence in decoded]

~~~~~ PREDICTIONS

[<tf.Tensor: shape=(), dtype=string, numpy=b'bin red at s nine
again'>]
```